

# Provable Infinite-Horizon Real-Time Planning for Repetitive Tasks

## Abstract

In manufacturing, robots often have to perform highly-repetitive manipulation tasks in structured environments. In this work we are interested in the settings where the tasks are similar, yet not identical (e.g., due to uncertain orientation of objects) and motion planning needs to be extremely fast. Preprocessing-based approaches prove to be very beneficial in these settings—they analyze the configuration-space offline to generate some auxiliary information which can then be used in the query phase to speedup planning times. Typically, the tighter the requirement is on query times the larger the memory footprint will be. In particular, for high-dimensional spaces, providing real-time planning capabilities is impractical. Moreover, as far as we are aware of, none of the general-purpose algorithms come with *provable* guarantees on the real-time performance. To this end, we propose a preprocessing-based method that provides provable bounds on the query time while incurring only a small amount of memory overhead in the query phase. We evaluate our method on a 7-DOF robot arm and show a speedup of over tenfold in query time when compared to the PRM algorithm, while guaranteeing a maximum query time of less than 4 milliseconds.

## 1 Introduction

We consider the problem of planning robot motions for highly-repetitive tasks while ensuring bounds on the planning times. As a running example, consider the problem of a robot picking up objects from a high-speed conveyor belt and placing them into bins (see Fig. 1). While the environment does not change, the start and goal in each task are similar, yet not identical to the start and goal in previous tasks. Differences in the exact start and goal positions may be due to uncertainty in the environment as objects may be placed on different parts of the conveyor belt or in different orientations. Fast planning times immediately correspond to faster unloading capabilities. Moreover, if the planner cannot *guarantee* to pick items from the conveyor in a timely manner, the system is required to account for missed items by e.g., additional conveyor belts that will redirect items back to the robot—a costly backup in terms of both time and space.

Copyright © 2019, Association for the Advancement of Artificial Intelligence (www.aaai.org). All rights reserved.



Figure 1: Motivating scenario—a robot (PR2) picking up objects from a conveyor belt.

Clearly, every time a task is presented to the robot, it can compute a desired path. However, this may incur large on-line planning times that may be unacceptable in many settings. Alternatively, we could attempt to precompute for each start and goal pair a robot path. However, as the set of possible start and goal locations may be large, caching pre-computed paths for all these queries in advance is unmanageable in high-dimensional configuration spaces. Thus, we need to balance memory constraints while providing provable real-time query times.

As we detail in Sec. 2, there has been intensive work for fast online planning (Lehner and Albu-Schaffer 2018) and for learning from experience in known environments (Phillips et al. 2012; Phillips et al. 2013; Berenson, Abbeel, and Goldberg 2012b; Coleman et al. 2015). Similarly, compressing precomputed data structures in the context of motion-planning is a well-studied problem with efficient algorithms (Salzman et al. 2014; Dobson and Bekris 2014). However, to the extent of our knowledge, there is no approach that can *provably guarantee* that a solution will be found to *any* query with bounds on planning time and using a low memory footprint.

Our key insight is that given any state  $s$ , we can efficiently compute a set of states for which a greedy search towards  $s$  is collision free. Importantly, the runtime-complexity of such a greedy search is bounded and there is no need to

perform computationally-complex collision-detection operations. This insight allows us to generate in an offline phase a small set of so-called “attractor vertices” together with a path between each attractor state and the goal. In the query phase, a path is generated by performing a greedy search to an attractor state followed by the precomputed path to the goal. We describe our approach in Sec. 3 and analyze it in Sec. 4

We evaluate our approach in Sec. 5 in simulation on the PR2 robot<sup>1</sup> (see Fig. 1). We demonstrate a speedup of over ten-fold in query time when compared to the PRM algorithm with a little memory footprint of 0.2 Mbs and while guaranteeing a maximal query time of less than 4 milliseconds.

## 2 Related work

A straightforward approach to efficiently preprocess a known environment is using the PRM algorithm (Kavraki et al. 1996) which generates a *roadmap*. Once a dense roadmap has been pre-computed, any query can be efficiently answered online by connecting the source and goal to the roadmap. Query times can be significantly sped up by further preprocessing the roadmaps using landmarks (Paden, Nager, and Frazzoli 2017). Unfortunately, there is no guarantee that a query can be connected to the roadmap as PRM only provides *asymptotic* guarantees (Kavraki, Kolountzakis, and Latombe 1998). Furthermore, this connecting phase requires running a collision-detection algorithm which is typically considered the computational bottleneck in many motion-planning algorithms (LaValle 2006).

Recently, Lehner and Albu-Schäffer (Lehner and Albu-Schaffer 2018) suggest the repetition roadmap to extend the PRM for the case of multiple highly-similar scenarios. While their approach exhibits significant speedup in computation time, it still suffers from the previously-mentioned shortcomings.

A complementary approach to aggressively preprocess a given scenario is by minimizing collision-detection time. However this requires designing robot-specific circuitry (Murray et al. 2016) or limiting the approach to standard manipulators (Yang et al. 2018).

An alternative approach to address our problem is to precompute a set of complete paths into a library and given a query, attempt to match complete paths from the library to the new query (Berenson, Abbeel, and Goldberg 2012a; Jetchev and Toussaint 2013). Using paths from previous search episodes (also known as using experience) has also been an active line of work (Phillips et al. 2012; Phillips et al. 2013; Berenson, Abbeel, and Goldberg 2012b; Coleman et al. 2015). Some of these methods have been integrated with sparse motion-planning roadmaps (see e.g., (Salzman et al. 2014; Dobson and Bekris 2014)) to reduce the memory footprint of the algorithm. Unfortunately, none of the mentioned algorithms provide any of the guarantees required by our applications.

<sup>1</sup><http://www.willowgarage.com/pages/pr2/overview>

Our work bears resemblance to previous work on subgoal graphs (Uras and Koenig 2017; Uras and Koenig 2018) and to real-time planning (Koenig and Likhachev 2006; Koenig and Sun 2009; Korf 1990). However, in the former, the entire configuration space is preprocessed in order to efficiently answer queries between *any* pair of states which deems it applicable only to low-dimensional spaces. **Also their *connect* step is expensive as it requires a Dijkstra’s like search to connect to the subgoal graph.** Similarly, in the latter, to provide guarantees on planning time the search only looks at a finite horizon and interleaves planning and execution.

Finally, our notion of attractor states is similar to control-based methods that ensure safe operation over local regions of the free configuration space (Conner, Rizzi, and Choset 2003; Conner, Choset, and Rizzi 2006). These regions are then used within a high-level motion planner to compute collision-free paths.

## 3 Algorithm Framework

In this section we describe our algorithmic framework. We start (Sec. 3.1) by formally defining our problem and continue (Sec. 3.2) by describing the key idea that enables our approach. We then proceed (Sec. 3.3) to detail our algorithm and conclude (Sec. 3.4)

### 3.1 Problem formulation and assumptions

Let  $\mathcal{X}$  be the configuration space of a robot operating in a static environment. We are given in advance a start configuration  $s_{\text{start}} \in \mathcal{X}$  and some goal region  $G \subset \mathcal{X}$ . In the query phase we are given multiple queries  $(s_{\text{start}}, s_{\text{goal}})$  where  $s_{\text{goal}} \in G$  and for each query, we need to compute a collision-free path connecting  $s_{\text{start}}$  to  $s_{\text{goal}}$ .

We discretize  $\mathcal{X}$  into a state lattice  $\mathcal{S}$  such that any state  $s \in \mathcal{S}$  is connected to a set of successors via a mapping  $\text{Succs}: \mathcal{S} \rightarrow 2^{\mathcal{S}}$  and set  $G_{\mathcal{S}} := \mathcal{S} \cap G$  to be the states that reside in the goal region. We make the following assumptions:

- A1**  $G_{\mathcal{S}}$  is a relatively small subset of  $\mathcal{S}$ . Namely, it is feasible to exhaustively iterate over all states in  $G_{\mathcal{S}}$ . However, storing a path from  $s_{\text{start}}$  to each state in  $G_{\mathcal{S}}$  is infeasible.
- A2** The planner has access to a heuristic function  $h: \mathcal{S} \times \mathcal{S} \rightarrow \mathbb{R}$  which can estimate the distance between any two states in  $G_{\mathcal{S}}$ . Moreover the heuristic function should be *weakly-monotonic*, meaning that  $\forall s_1, s_2 \in G_{\mathcal{S}}$  where  $s_1 \neq s_2$ , it holds that,

$$h(s_1, s_2) \geq \min_{s'_1 \in \text{Succs}(s_1)} h(s'_1, s_2).$$

Namely, for any distinct pair of states  $(s_1, s_2)$  in  $G_{\mathcal{S}}$ , at least one of  $s_1$ ’s successors (also belonging to  $G_{\mathcal{S}}$ ) must have a heuristic value less than or equal to its heuristic value. Note that this assumption does not imply that  $G$  is entirely collision free.

**A3** The planner has access to a tie breaking rule that can be used to define a total order<sup>2</sup> over all states with the same heuristic value to a given state.

These assumptions allow us to establish strong theoretical properties regarding the efficiency of our planner. Namely, that within a known bounded time, we can compute a collision-free path from  $s_{\text{start}}$  to any state in  $G_S$ .

### 3.2 Key idea

Add a figure that demonstrates the complete approach.

Our algorithm relies heavily on the notion of a greedy search. Thus, before we provide an overview of our algorithm, we formally define the notions of greedy successor and greedy search.

**Definition 1.** Let  $s$  be some state and  $h(\cdot)$  be some heuristic function. A state  $s' \in \text{Succ}(s)$  is said to be a greedy successor of  $s$  according to  $h$  if it has the minimal  $h$ -value among all of  $s$ 's successors.

Note that if  $h$  is weakly monotone (Assumption **A2**) and we have some tie-breaking rule (Assumption **A3**), then the greedy successor of a state is unique. In the rest of the text, when we use the term greedy successor, we assume that it is unique.

**Definition 2.** Given a heuristic function  $h(\cdot)$ , an algorithm is said to be a greedy search if for every state it returns its greedy successor.

Our key insight is to precompute in an offline phase regions where a greedy search to a certain state is guaranteed to be collision free and use these regions in the query phase. Specifically, in the preprocessing phase,  $G_S$  is decomposed into two finite sets of (possibly overlapping) subregions  $\mathcal{R}$  and  $\hat{\mathcal{R}}$ . Subregions in  $\hat{\mathcal{R}}$  only contain states that are in collision. Each subregion  $R_i \in \mathcal{R}$  is a hyper-ball defined using a center which we refer to as the ‘‘attractor state’’  $s_i^{\text{attractor}}$  and a radius  $r_i$ . These regions are constructed in such a way that the following two properties hold

- P1** For any goal state  $s_{\text{goal}} \in R_i \cap G_S$ , a greedy search with respect to  $h(s, s_i^{\text{attractor}})$  over  $\mathcal{S}$  starting at  $s_{\text{goal}}$  will result in a collision-free path to  $s_i^{\text{attractor}}$ .
- P2** The union of all the subregions completely cover  $G_S$ . Namely,  $\forall s \in G_S, \exists R \in \mathcal{R} \cup \hat{\mathcal{R}} \text{ s.t. } s \in R$ .

In the preprocessing stage, we precompute a library of collision-free paths  $\mathcal{L}$  which includes one path from  $s_{\text{start}}$  to each attractor state. In the query phase, given a query  $s_{\text{goal}}$ , we (i) identify a region  $R_i$  such  $s_{\text{goal}} \in R_i$  (using the pre-computed radii  $r_i$ ), (ii) run a greedy search towards  $s_i^{\text{attractor}}$  by greedily choosing at every point the successor that minimizes  $h$  and (iii) append this path with the precomputed path in  $\mathcal{L}$  to  $s_{\text{start}}$  to obtain the complete plan.

<sup>2</sup>A total order is a binary relation on some set which is antisymmetric, transitive, and a connex relation.




Figure 2: Visualization of Alg 2. Subregion  $R_i$  (green) grown from  $s_i^{\text{attractor}}$  in a goal region  $G_S$  (grey) containing an obstacle (red). Frontier states and first state not in  $R_i$  are depicted by circles and a cross, respectively.

### 3.3 Algorithm

**Preprocessing Phase** The preprocessing phase of our algorithm, detailed in Alg. 1, takes as input the start state  $s_{\text{start}}$  and the goal region  $G_S$  and outputs a set of subregions  $\mathcal{R}$  and the corresponding library of paths  $\mathcal{L}$  from each  $s_i^{\text{attractor}}$  to  $s_{\text{start}}$ .

The algorithm covers  $G_S$  by iteratively finding a state  $s$  not covered<sup>3</sup> by any region and computing a new region centered at  $s$ . To ensure that  $G_S$  is complete covered (Property **P2**) we maintain a set  $V$  of valid (collision free) and a set  $I$  of invalid (in collision) states called *frontier states* (lines 3 and 4, respectively). We start by initializing  $V$  with some random state in  $G_S$  and iterate until both  $V$  and  $I$  are empty, which will ensure that  $G_S$  is indeed covered.

At every iteration, we pop a state from  $V$  (line 8), and if there is no region covering it, we add it as a new attractor state and compute a path  $\pi_i$  to  $s_{\text{start}}$  (line 11). We then compute the corresponding region (line 12 and Alg. 2).

As we will see shortly, computing a region corresponds to a Dijkstra-like search centered at the attractor state. The search terminates with the region’s radius  $r_i$  and a list of frontier states that comprise of the region’s boundary. The valid and invalid states are then added to  $V$  and  $I$ , respectively (lines 13 and 14).

Once  $V$  gets empty the algorithm starts to search for states which are valid and yet uncovered by growing regions around the states popped from  $I$  (lines 16-20). If a valid and uncovered state is found, it is added to  $V$  and the algorithm goes back to computing subregions (lines 21-23), otherwise if  $I$  also gets empty, the algorithm terminates and it is guaranteed that each valid state contained in  $G_S$  is covered under at least one subregion.

**Reachability Search** The core of our planner lies in the way we compute the subregions (Alg. 2 and Fig. 2) which we call a ‘‘Reachability Search’’. The algorithm maintains a set of *reachable* states  $S_{\text{reachable}}$  for which property **P1** holds. As we will see this will ensure that in the query

<sup>3</sup>Here, a state  $s$  is said to be covered if there exists some region  $R \in \mathcal{R}$  such that  $s \in R$ .

---

**Algorithm 1** Goal Region Preprocessing

---

**Inputs:**  $G_S, s_{\text{start}}$   $\triangleright$  goal region and start state  
**Outputs:**  $\mathcal{R}, \mathcal{L}$   $\triangleright$  subregions and corresponding paths to  $s_{\text{start}}$

```

1: procedure PREPROCESSREGION( $G_S$ )
2:    $s \leftarrow \text{SAMPLEVALIDSTATE}(G_S)$ 
3:    $V \leftarrow \{s\}$   $\triangleright$  valid frontier states initialized to a random state
4:    $I = \emptyset$   $\triangleright$  invalid frontier states
5:    $i \leftarrow 0$   $\mathcal{L} = \emptyset$   $\mathcal{R} = \emptyset$   $\hat{\mathcal{R}} = \emptyset$ 
6:   while  $V$  and  $I$  are not empty do
7:     while  $V$  is not empty do
8:        $s \leftarrow V.\text{pop}()$ 
9:       if  $\nexists R \in \mathcal{R}$  s.t.  $s \in R$  then  $\triangleright s$  is not covered
10:         $s_i^{\text{attractor}} \leftarrow s$ 
11:         $\pi_i = \text{PLANPATH}(s_i^{\text{attractor}}, s_{\text{start}})$ ;  $\mathcal{L} \leftarrow \mathcal{L} \cup \{\pi_i\}$ 
12:         $(\text{OPEN}, r_i) \leftarrow \text{COMPUTEREACHABILITY}(s_i^{\text{attractor}})$ 
13:        insert Valid( $\text{OPEN}$ ) in  $V$ 
14:        insert Invalid( $\text{OPEN}$ ) in  $I$ 
15:         $R_i \leftarrow (s_i^{\text{attractor}}, r_i)$ ;  $i \leftarrow i + 1$   $\mathcal{R} \leftarrow \mathcal{R} \cup \{R_i\}$ 
16:     while  $I$  is not empty do
17:        $s \leftarrow I.\text{pop}()$ 
18:       if  $\nexists R \in \mathcal{R} \cup \hat{\mathcal{R}}$  s.t.  $s \in R$  then  $\triangleright s$  is not covered
19:         $(X, r) \leftarrow \text{SEARCHVALIDUNCOVEREDSTATES}(s)$ 
20:         $\hat{R} \leftarrow (s, r)$ ;  $\hat{\mathcal{R}} \leftarrow \hat{\mathcal{R}} \cup \{\hat{R}\}$   $\triangleright$  invalid region
21:        if  $X$  is not empty then  $\triangleright$  no valid state found
22:          insert  $X$  in  $V$ 
23:        break
24:   return  $\mathcal{R}, \mathcal{L}$ 

```

---

phase, we can run a greedy search from any reachable state  $s \in S_{\text{reachable}}$  and it will terminate in the attractor state. The following definition formally captures the notion of a reachable state.

**Definition 3.** Given some attractor state  $s_i^{\text{attractor}}$ , a state  $s$  is said to be reachable with respect to  $s_i^{\text{attractor}}$  if either (i)  $s = s_i^{\text{attractor}}$  or (ii) the greedy successor of  $s$  is reachable respect to  $s_i^{\text{attractor}}$ .

The algorithm computes a subregion that covers the maximum number of reachable states that can fit into a hyper-ball defined by  $h(s, s_i^{\text{attractor}})$ . The search maintains a priority queue OPEN ordered according to  $h(s, s_i^{\text{attractor}})$ . Initially, the predecessors of  $s_i^{\text{attractor}}$  are inserted in the OPEN (line 3). For each expanded predecessor, if its valid greedy successor is in  $S_{\text{reachable}}$ , then the predecessor is also labeled as reachable (lines 10 and 11).

The algorithm terminates when the search pops a state which is valid but does not have a greedy successor state in  $S_{\text{reachable}}$  (line 12). Intuitively, this corresponds to the condition when the reachability search exists an obstacle (see Fig. 2). At termination, all the states within the boundary of radius  $r_i$  (excluding the boundary) are reachable.

**Query Phase** Given a query goal state  $s_{\text{goal}} \in G_S$  our algorithm, detailed in Alg. 3, starts by finding a subregion  $R_i \in \mathcal{R}$  which covers it (Line. 9). Namely, a region  $R_i$  for which  $h(s_{\text{goal}}, s_i^{\text{attractor}}) < r_i$ . We then run a greedy search

---

**Algorithm 2** Reachability Search

---

```

1: procedure COMPUTEREACHABILITY( $s_i^{\text{attractor}}$ )
2:    $S_{\text{reachable}} \leftarrow \{s_i^{\text{attractor}}\}$   $\triangleright$  Reachable set
3:    $\text{OPEN} \leftarrow \{\text{Preds}(s_i^{\text{attractor}})\}$   $\triangleright$  key:  $h(s, s_i^{\text{attractor}})$ 
4:    $\text{CLOSED} \leftarrow \emptyset$ 
5:    $r_i \leftarrow 0$ 
6:   loop
7:      $s \leftarrow \text{OPEN}.\text{pop}()$ 
8:     insert  $s$  in  $\text{CLOSED}$ 
9:      $s'_g \leftarrow \arg \min_{s' \in \text{Succ}(s)} h(s', s_i^{\text{attractor}})$   $\triangleright$  greedy successor
10:    if  $s'_g \in S_{\text{reachable}}$  and Valid(edge( $s, s'_g$ )) then
11:       $S_{\text{reachable}} \leftarrow S_{\text{reachable}} \cup \{s\}$   $\triangleright s$  is greedy
12:    else if Valid( $s$ ) then
13:      return  $r_i$ 
14:     $r_i \leftarrow h(s, s_i^{\text{attractor}})$ 
15:    for each  $p \in \text{Preds}(s)$  do
16:      if  $p \notin \text{CLOSED}$  then
17:        insert  $p$  in  $\text{OPEN}$  with priority  $h(p, s_i^{\text{attractor}})$ 

```

---



---

**Algorithm 3** Query

---

```

1: procedure FINDGREEDYPATH( $s_1, s_2$ )
2:    $s_{\text{curr}} \leftarrow s_1$ ;  $\pi \leftarrow \emptyset$ 
3:   while  $s_{\text{curr}} \neq s_2$  do
4:      $\pi \leftarrow \pi \cdot s_{\text{curr}}$   $\triangleright$  append current state to path
5:      $s_{\text{curr}} \leftarrow \arg \min_{s \in \text{Succ}(s_{\text{curr}})} h(s, s_2)$   $\triangleright$  greedy successor
6:   return  $\pi$ 
7: procedure COMPUTE PATH( $s_{\text{goal}}$ )
8:   for each  $R_i \in \mathcal{R}$  do
9:     if  $h(s_{\text{goal}}, s_i^{\text{attractor}}) < r_i$  then  $\triangleright R_i$  covers  $s_{\text{goal}}$ 
10:       $\pi_g \leftarrow \text{FINDGREEDYPATH}(s_{\text{goal}}, s_i^{\text{attractor}})$ 
11:      return  $\pi_g \cdot \pi_i$   $\triangleright$  append  $\pi_g$  to  $\pi_i \in \mathcal{L}$ 

```

---

starting from  $s_{\text{goal}}$  by iteratively finding for each state  $s$  the successor with the minimum heuristic  $h(s, s_i^{\text{attractor}})$  value until the search reaches  $s_i^{\text{attractor}}$  (Lines 10 and 1- 6). The greedy path  $\pi_g$  is then appended to the corresponding pre-computed path  $\pi_i \in \mathcal{L}$  (Line 11). Note that at no point do we need to perform collision checking in the query phase.

### 3.4 Implementation details

Recall that in the query phase we iterate over all regions to find a region that covers  $s_{\text{goal}}$ . In the worst case we will have to go over all regions. However, as we will see in our evaluation (Sec. 5), our algorithm typically covers most of the goal region  $G_S$  using a few very large regions (namely, with a large radii  $r_i$ ) and the rest of  $G_S$  is covered by a large number of very small regions.

Thus, if we order our regions according to their corresponding radii, there is a higher chance of finding a covering region faster. While this optimization does not change our worst-case analysis (Sec. 4), we observe that in practice (Sec. 5) it can dramatically speed up our query times.



## 4 Analysis

In this section we formally prove that our algorithm is correct (Sec. 4.1) and analyze its computational complexity (Sec. 4.2).

### 4.1 Correctness

To prove that our algorithm is correct, we show that indeed all states of every region are reachable and we can identify if a state belongs to a region using its associated radius. Furthermore, we show that a path obtained by a greedy search within any region is valid and that all states in  $G_S$  are covered by some region. These notions are captured by the following set of lemmas.

**Lemma 1.** *Let  $S_{\text{reachable}}$  be the set of states computed by Alg. 2 for some attractor vertex  $s_i^{\text{attractor}}$ . Every state  $s \in S_{\text{reachable}}$  is reachable with respect to  $s_i^{\text{attractor}}$ .*

*Proof.* The proof is constructed by an induction over the states added to  $S_{\text{reachable}}$ . The base of the induction is trivial as the first state added to  $S_{\text{reachable}}$  is  $s_i^{\text{attractor}}$  (Line 2) which, by definition, is reachable with respect to  $s_i^{\text{attractor}}$ . A state  $s$  is added to  $S_{\text{reachable}}$  only if its greedy successor is in  $S_{\text{reachable}}$  (Line 9) which by the induction hypothesis is reachable with respect to  $s_i^{\text{attractor}}$ . This implies by definition that  $s$  is reachable with respect to  $s_i^{\text{attractor}}$ . Note that this argument is true because the greedy successor of every state is unique (Assumption A3).  $\square$

**Lemma 2.** *Let  $S_{\text{reachable}}$  be the set of states computed by Alg. 2 for some attractor vertex  $s_i^{\text{attractor}}$ . A state  $s$  is in  $S_{\text{reachable}}$  iff  $h(s, s_i^{\text{attractor}}) < r_i$ .*

*Proof.* Alg. 2 orders the nodes to be inserted to  $S_{\text{reachable}}$  according to  $h(s, s_i^{\text{attractor}})$  (Line 3). As our heuristic function is weakly monotonic (Assumption A2), the value of  $r_i$  monotonically increases as the algorithm adds states to  $S_{\text{reachable}}$  (Line 14). Thus, for every state  $s \in S_{\text{reachable}}$ , we have that  $h(s, s_i^{\text{attractor}}) < r_i$ .

For the opposite direction, assume that there exists a state  $s \notin S_{\text{reachable}}$  such that  $h(s, s_i^{\text{attractor}}) \leq r_i$ . As  $s \notin S_{\text{reachable}}$ , Alg. 2 terminated due to a node  $s'$  that was popped from the open list with  $h(s', s_i^{\text{attractor}}) \geq r_i$ . Again, using the fact that our heuristic function is weakly monotonic (Assumption A2) we get a contradiction to the fact that  $h(s, s_i^{\text{attractor}}) \leq r_i$ .  $\square$

**Lemma 3.** *Let  $R_i \in \mathcal{R}$  be a region computed by Alg. 2 for some attractor vertex  $s_i^{\text{attractor}}$ . A greedy search with respect to  $h(s, s_i^{\text{attractor}})$  starting from any valid state  $s \in R_i$  is complete and valid.*

*Proof.* Given a state  $s \in R_i$ , we know that  $s \in S_{\text{reachable}}$  (Lemma 2) and that it is reachable with respect to  $s_i^{\text{attractor}}$  (Lemma 1). It is easy to show (by induction) that any greedy search starting at a state  $s \in S_{\text{reachable}}$  will only output states in  $S_{\text{reachable}}$ . Furthermore, a state is added to  $S_{\text{reachable}}$

	PRM (T)	PRM (2T)	PRM (4T)	Our method
Planning time [ms]	20.2 (28.2)	21.8 (34.1)	22.9 (31.9)	0.48 (1.01)
Success rate [%]	92	97	100	100

Table 1: Experimental results comparing our method with PRM. The table shows the mean/worst-case planning times and success rates for our method and for PRM preprocessed with equal, double and quadruple the time that our method takes in precomputation (T = 984 seconds).

only if the edge connecting to its greedy successor is valid (Line 10). Thus, if  $s \in S_{\text{reachable}}$  is valid, the greedy search with respect to  $h(s, s_i^{\text{attractor}})$  starting from  $s$  is complete and valid.  $\square$

**Lemma 4.** *At the end of Alg. 1, every state  $s \in G_S$  is covered by some region  $R \in \mathcal{R}$ .*

*Proof.* Assume that this does not hold and let  $s \in G_S$  be a state that is not covered by any region but has a neighbor that is covered. If  $s$  is valid, then it would have been in the valid frontier states  $V$  and either been picked to be an attractor state (Line 10) or covered by an existing region (Alg. 2). A similar argument holds if  $s$  is not valid.  $\square$

From the above we can immediately deduce the following corollary:

**Corollary 1.** *After preprocessing the goal region  $G_S$  (Alg. 1 and 2), in the query phase we can compute a valid path for any valid state  $s \in G_S$  using Alg. 3.*

### 4.2 Time Complexity of Query Phase

[Do we want to do query complexity of preprocessing phase?.](#)

The query time comprises of (i) finding the containing subregion  $R_i$  and (ii) running the greedy search to  $s_i^{\text{attractor}}$ . Step (i) requires iterating over all subregions (in the worst case) which takes  $O(|\mathcal{R}|)$  steps while step (ii) requires expanding the states along the path from  $s_{\text{goal}}$  to  $s_i^{\text{attractor}}$  which requires  $O(\mathcal{D})$  expansions where  $\mathcal{D}$  is the depth of the deepest subregion. For each expansion we need to find the greedy successor, considering at most  $b$  successor, where  $B$  is the maximal branching factor of our graph. We can measure the depth of each subregion in Alg. 2 by keeping track of the depth of each expanded state from the root i.e. the  $s_i^{\text{attractor}}$ . Hence, overall the query phases takes  $O(|\mathcal{R}| + \mathcal{D} \cdot b)$  operations. The maximal query time can also be empirically computed after the preprocessing phase.

**Remark:** Note that we can also bound the number of expansions required for the query phase by bounding the maximum depth of the subregions. We can do that by terminating Alg. 2 when the  $R_i$  reaches the maximum depth or if the existing termination condition (line 12) is satisfied. Having said that, this may come at the price of increasing the number of regions.

## 5 Evaluation

We evaluated our algorithm by getting some preliminary results on the PR2 robot for the single-arm (7 DOF) planning problem. The task here is to repeatedly pick up objects from a conveyor belt and put them in a bin. We define the task-relevant goal region  $G$  by bounding the position and orientation of the end effector.

We compared our approach with the PRM algorithm in terms of success rate and planning times (see Table 1) for 100 uniformly sampled goal states from  $G$ . Preprocessing (Alg. 1) took 984 seconds and returned 1,865 subregions. For a fair comparison, for PRM the paths from all the nodes in  $G$  to  $s_{\text{start}}$  were precomputed. In the query phase if PRM's connect operation fails for a given query, we consider it a failure. We also bootstrap PRM with 20 goal states from  $G$ . Note that the worst-case time for our method shown in these results ( $\sim 1$  millisecond) is the empirical one and not the computed provable time bound which is 4 milliseconds for this environment.

## 6 Conclusion

We proposed a preprocessing-based motion planning algorithm that provides provable real-time performance guarantees for repetitive tasks and showed preliminary results. We aim to perform experiments on the hardware with the conveyor setup to demonstrate the efficiency of our planner on a real world system. Moreover on the theoretical side, we aim to provide guarantees on the solution quality.

Talk about picking the right attractor points to minimize the number of regions and in general minimizing the overall query time.

## References

- [Berenson, Abbeel, and Goldberg 2012a] Berenson, D.; Abbeel, P.; and Goldberg, K. 2012a. A robot path planning framework that learns from experience. In *IEEE International Conference on Robotics and Automation (ICRA)*, 3671–3678.
- [Berenson, Abbeel, and Goldberg 2012b] Berenson, D.; Abbeel, P.; and Goldberg, K. 2012b. A robot path planning framework that learns from experience. In *IEEE International Conference on Robotics and Automation (ICRA)*, 3671–3678.
- [Coleman et al. 2015] Coleman, D.; Şucan, I. A.; Moll, M.; Okada, K.; and Correll, N. 2015. Experience-based planning with sparse roadmap spanners. In *IEEE International Conference on Robotics and Automation (ICRA)*, 900–905.
- [Conner, Choset, and Rizzi 2006] Conner, D. C.; Choset, H.; and Rizzi, A. 2006. Integrated planning and control for convex-bodied nonholonomic systems using local feedback control policies. Technical Report CMU-RI-TR-06-34, Carnegie Mellon University, Pittsburgh, PA.
- [Conner, Rizzi, and Choset 2003] Conner, D. C.; Rizzi, A.; and Choset, H. 2003. Composition of local potential functions for global robot control and navigation. In *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, volume 4, 3546–3551.
- [Dobson and Bekris 2014] Dobson, A., and Bekris, K. E. 2014. Sparse roadmap spanners for asymptotically near-optimal motion planning. *IJRR* 33(1):18–47.
- [Jetchev and Toussaint 2013] Jetchev, N., and Toussaint, M. 2013. Fast motion planning from experience: trajectory prediction for speeding up movement generation. *Autonomous Robots* 34(1-2):111–127.
- [Kavraki et al. 1996] Kavraki, L. E.; Švestka, P.; Latombe, J.-C.; and Overmars, M. H. 1996. Probabilistic roadmaps for path planning in high-dimensional configuration spaces. *IEEE Trans. Robotics and Automation* 12(4):566–580.
- [Kavraki, Kolountzakis, and Latombe 1998] Kavraki, L. E.; Kolountzakis, M. N.; and Latombe, J. 1998. Analysis of probabilistic roadmaps for path planning. *IEEE Trans. Robotics and Automation* 14(1):166–171.
- [Koenig and Likhachev 2006] Koenig, S., and Likhachev, M. 2006. Real-time adaptive A\*. In *International joint conference on Autonomous agents and multiagent systems*, 281–288. ACM.
- [Koenig and Sun 2009] Koenig, S., and Sun, X. 2009. Comparing real-time and incremental heuristic search for real-time situated agents. *Autonomous Agents and Multi-Agent Systems* 18(3):313–341.
- [Korf 1990] Korf, R. E. 1990. Real-time heuristic search. *Artificial Intelligence* 42(2-3):189–211.
- [LaValle 2006] LaValle, S. M. 2006. *Planning Algorithms*. Cambridge University Press.
- [Lehner and Albu-Schaffer 2018] Lehner, P., and Albu-Schaffer, A. 2018. The Repetition Roadmap for Repetitive Constrained Motion Planning. *IEEE Robotics and Automation Letters*. to appear.
- [Murray et al. 2016] Murray, S.; Floyd-Jones, W.; Qi, Y.; Sorin, D. J.; and Konidaris, G. 2016. Robot Motion Planning on a Chip. In *Robotics: Science and Systems (RSS)*.
- [Paden, Nager, and Frazzoli 2017] Paden, B.; Nager, Y.; and Frazzoli, E. 2017. Landmark guided probabilistic roadmap queries. In *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, 4828–4834.
- [Phillips et al. 2012] Phillips, M.; Cohen, B. J.; Chitta, S.; and Likhachev, M. 2012. E-graphs: Bootstrapping planning with experience graphs. In *Robotics: Science and Systems (RSS)*.
- [Phillips et al. 2013] Phillips, M.; Dornbush, A.; Chitta, S.; and Likhachev, M. 2013. Anytime incremental planning with e-graphs. In *IEEE International Conference on Robotics and Automation (ICRA)*, 2444–2451.

- [Salzman et al. 2014] Salzman, O.; Shaharabani, D.; Agarwal, P. K.; and Halperin, D. 2014. Sparsification of motion-planning roadmaps by edge contraction. *IJRR* 33(14):1711–1725.
- [Uras and Koenig 2017] Uras, T., and Koenig, S. 2017. Feasibility Study: Subgoal Graphs on State Lattices. In *Symposium on Combinatorial Search, SOCS*, 100–108.
- [Uras and Koenig 2018] Uras, T., and Koenig, S. 2018. Fast Near-Optimal Path Planning on State Lattices with Subgoal Graphs. In *Symposium on Combinatorial Search, SOCS*, 106–114.
- [Yang et al. 2018] Yang, Y.; Merkt, W.; Ivan, V.; Li, Z.; and Vijayakumar, S. 2018. HDRM: A Resolution Complete Dynamic Roadmap for Real-Time Motion Planning in Complex Scenes. *IEEE Robotics and Automation Letters* 3(1):551–558.